

Charging Infrastructure for Electric Vehicles: Challenges, AI/ML-Enabled Solutions, and Future Directions

Fulpagare Pradnyesh Sharad

Department of Computer Science and Application
MIT World Peace University
Pune Maharashtra
1272250462@mitwpu.edu.in

Nangare Ishaa Vivek

Department of Computer Science and Application
MIT World Peace University
Pune Maharashtra
1272250455@mitwpu.edu.in

Abstract—Basic This paper reviews the state of electric-vehicle (EV) charging infrastructure, identifies both technical and non-technical challenges, and proposes AI/ML-based solutions for smarter and more reliable charging. Key focus areas include grid impacts, station reliability, interoperability, and site-planning issues. To address these, three AI/ML case studies are presented: (1) short-term charging demand forecasting, (2) reinforcement-learning (RL) based smart-charging and scheduling, and (3) predictive maintenance for charging stations. Datasets and simulation tools are explored, baseline algorithms are implemented, and evaluation is conducted using forecasting and system-level metrics. The paper concludes with a discussion on policy, privacy, and future research directions such as vehicle-to-grid (V2G) integration and federated learning.

Key factual context: global public charger growth and the availability of public charging datasets are changing rapidly and motivate urgent work on reliability and smart control.

1. Introduction

This Wide EV adoption depends critically on charging infrastructure quality: coverage, uptime, affordability, and grid compatibility. Recent rapid additions of public chargers worldwide underscore the importance of robust operations and smart controls; yet reliability and equity gaps persist, and the grid faces growing stress from unscheduled high-power charging. AI/ML offers tools for forecasting, scheduling, and health-monitoring, but practical deployment faces data fragmentation, privacy concerns, and heterogeneity of hardware and protocols. Contributions (this paper):

1. A novel hybrid Spatio-Temporal Graph LSTM (ST-G-LSTM) architecture for station-level short-term demand forecasting (real-time, interpretable).
2. A Constrained Multi-Agent Reinforcement Learning (C-MARL) scheduler: fairness-aware, grid-aware allocation with provable constraint satisfaction via Lagrangian duals.

3. A Federated Predictive Maintenance pipeline combining unsupervised autoencoder anomaly detection + local causal attribution (SHAP-like) for actionable maintenance recommendations while preserving privacy.
4. Reproducible experimental protocol (datasets, synthetic scenario generator, evaluation metrics) to validate performance under realistic grid/traffic conditions.

2. Literature review

- Global charging trends and scale: IEA Global EV Outlook 2025 reports rapid growth in public charging infrastructure (over ~1.3M public chargers added in 2024). This motivates urgent scalability and reliability research.
- Public datasets: ACN-Data (Caltech) provides session-level workplace charging logs used widely for demand modelling.
- Reinforcement Learning for EV charging and smart scheduling is an active area with dedicated reviews summarizing RL successes and limitations for load-flattening and grid services.
- Vehicle-to-Grid (V2G) integration and grid services: recent technical reports (e.g., US DOE VGI assessment) emphasize V2G potential and integration challenges.
- Reliability studies: investigations into real-world uptime show significant gaps between reported vs observed availability, reinforcing the need for predictive maintenance and better monitoring.

Gap identified: existing ML solutions often target single problems (forecast or scheduling or maintenance)

in isolation and rely on centralized, privacy-invasive datasets. Our work integrates forecasting, constrained scheduling, and federated maintenance with explicit deployment guidance.

3. Problem formulation and system architecture

3.1 System actors and data flow

- **EV users:** session arrival time, desired departure, requested energy, battery state-of-charge (SoC) if available.
- **Charging stations / EVSE units:** charger power rating, number of ports, sensor telemetry (temp, current, error codes), local controller.
- **Distribution grid node:** transformer capacity, current load, dynamic tariffs or node-level constraints.
- **Operator / Fleet manager:** receives forecasts, issues schedule policies, deploys maintenance jobs.

3.2 Objectives (multi-objective)

- Minimize user waiting time and unmet energy requests.
- Minimize peak power draw and grid constraint violations.
- Maximize station uptime by early fault detection.
- Respect fairness across users and sites (avoid repeated preemption).

4. Case Study A Short-term station-level demand forecasting

4.1 Motivation

Accurate 15-min to 24-hour forecasts at station granularity enable proactive resource allocation, V2G scheduling, and grid-aware control.

4.2 Proposed model: ST-G-LSTM (novel)

A hybrid architecture combining:

- **Spatial layer:** Graph Neural Network (GNN) over charging-station graph (edges = road connectivity / electrical proximity / shared transformers) to capture spatial correlation.
- **Temporal layer:** Stacked LSTM with attention across time to capture daily/weekly seasonality and event effects.
- **Exogenous inputs:** weather, calendar (holidays), local events, real-time traffic density.
- **Probabilistic head:** Quantile forecasting (e.g., QR loss) to provide prediction intervals for operational risk management.

Model details (brief):

- Input per node $x_t^{(i)} = [\text{recent sessions, occupancy, hour-of-day, temp, } \dots]$.
- Spatial aggregation: $\tilde{x}_t^{(i)} = \sum_{j \in N(i)} w_{ij} \cdot \text{MLP}(x_t^{(j)})$ where w_{ij} learned via attention.
- LSTM processes $\tilde{x}_{t-k:t}^{(i)}$; attention over past K windows; quantile output layer.

4.3 Training & evaluation

- Loss: pinball/quantile losses + temporal smoothness regularizer.
- Baselines: ARIMA, Random Forest, single-node LSTM.
- Metrics: MAE, RMSE, SMAPE, prediction interval coverage probability (PICP).
- Data: ACN-Data plus synthetic augmentation (described in Section 6).

4.4 Novelty & expected benefits

- The spatial + probabilistic approach gives calibrated uncertainty that operators can use to set conservative power reservations and avoid overloads.

5. Case Study B Constrained Multi-Agent Reinforcement Learning (C-MARL) for Smart Charging

5.1 Problem as constrained MDP

Each charger agent decides real-time charging rates $a_t \in [0, P_{max}]$ for connected EVs. Global constraints: transformer capacity C_t and grid node power contracts. Reward mixes user service quality and grid cost; constraints enforce hard safety via penalties and Lagrangian multipliers.

5.2 Algorithm (novel combination)

- **Actor-Critic MARL:** Each station runs a local actor; a central critic uses partial observations to compute value and constraint gradients.
- **Lagrangian constraint module:** global constraints handled with learned multipliers updated via stochastic gradient ascent (dual updates) to ensure constraint satisfaction in expectation.
- **Fairness layer:** per-user debt counters and softmax allocation to avoid repeated preemption.
- **Robustness:** domain randomization for grid conditions during training and worst-case evaluation.

5.3 Implementation notes

- Use centralized training with decentralized execution (CTDE).

- Observation includes local SoC, predicted arrivals (from ST-G-LSTM), local queue, and a compact global congestion signal.
- Baselines: naive greedy, earliest-deadline-first (EDF), convex optimization schedulers (non-learning).

5.4 Expected outcome

- Compared to baselines, C-MARL should reduce peak power (kW) and waiting times while honoring transformer constraints; fairness indices improve compared to greedy.

6. Case Study C — Federated Predictive Maintenance with Causal Explanations

6.1 Motivation

Charging stations show real-world low uptime and inconsistent status reporting; predictive maintenance can reduce downtime, but raw telemetry is privacy-sensitive. A distributed, privacy-preserving approach is required.

6.2 Pipeline (novel)

1. Local preprocessing & feature extraction: each EVSE extracts standardized features (current harmonics, connector insertion events, temperature trends, error code histograms).
2. Federated unsupervised anomaly detection: clients train local autoencoders (or temporal convolutional variational autoencoders, TCVAE) and share encrypted model updates to an aggregator. Secure aggregation + differential privacy (DP) ensure privacy.
3. Local causal explanation: when anomaly score exceeds threshold, a compact local causal explainer (lightweight SHAP surrogate or perturbation-based causal probe) runs to identify likely failure causes (e.g., overheating, connector wear, firmware faults).
4. Operator dashboard: aggregated anonymized health map and prioritized maintenance tickets.

6.3 Algorithms & novelty

- Use Federated TCVAE with periodic personalized fine-tuning on-client to capture local hardware idiosyncrasies.
- Introduce explanation-aware thresholds: thresholding uses both anomaly score and causal-explanation confidence to reduce false positives.
- Incorporate temporal precursors to predict faults 24–72 hours ahead rather than just detect anomalies at failure time.

6.4 Data & evaluation

- Data sources: simulated sensor streams (see Section 7), any available operator logs (anonymized), and benchmark small datasets from manufacturers if accessible.

- Metrics: precision/recall for failure prediction, lead time before failure, reduction in mean time to repair (MTTR) in simulation.

Privacy & security: the pipeline integrates secure aggregation and DP; we prescribe parameter settings for epsilon to balance privacy/utility (detailed in Appendix).

7. Datasets, simulation environment and reproducible protocol

7.1 Public datasets

- ACN-Data: workplace charging sessions.
- Grid models / transformer capacities: public distribution network templates (OpenDSS feeders) and synthetic feeder models for scenario testing.
- Event & traffic data: municipal open traffic APIs or synthetic Poisson arrivals scaled to EV penetration.

7.2 Synthetic scenario generator (novel, shared artifact)

We provide an open synthetic generator to produce joint mobility + charging + grid scenarios:

- Road-to-station assignment with probabilistic travel time.
- Arrival models: mixture of Poisson + daily seasonality + event spikes.
- Charger fault injection module: controlled fault seeds (thermal rise, connector wear) to test predictive maintenance.
- Transformer and feeder constraints with stochastic load background profiles.

A public GitHub repo (suggested) will include generator code, configuration YAMLS for scenario presets (urban commuter, highway corridor, mixed residential), and seedable RNG for reproducibility.

7.3 Simulation stack (recommended)

- Charging / mobility: SUMO or custom event simulator.
- Grid power: OpenDSS or MATPOWER (python bindings).
- ML stack: PyTorch + PyG (for GNN), RLlib or Stable Baselines3 for MARL experiments.
- Federated learning: Flower or TensorFlow Federated for experimentation.

8. Experimental design & evaluation metrics

8.1 Forecasting experiments

- Horizon: 15min, 1h, 24h.
- Metrics: MAE, RMSE, SMAPE, PICP, CRPS (for probabilistic models).
- Ablation: spatial vs non-spatial, quantile vs point forecasts.

8.2 Scheduling experiments

- Scenarios: varying EV penetration (5%, 15%, 35% of vehicles EV), different transformer headrooms.
- Metrics: peak reduction (%), average wait time, % of unmet SOC requests, fairness index (Jain's index on per-user service).
- Robustness tests: sudden arrival surges, partial observability, dropped telemetry.

8.3 Maintenance experiments

- Inject faults at controlled times; measure detection lead time, precision/recall, and MTTR improvement under different operational policies (preventive vs reactive).

9. Deployment & policy considerations

9.1 Interoperability & standards

- Emphasize support for OCPP as baseline protocol and advocate for standardized telemetry schemas to enable federated ML.

9.2 Privacy & governance

- Recommend federated learning plus secure aggregation and DP for model updates. Include a stakeholder governance model where maintenance vendors can request anonymized aggregated evidence before dispatch.

9.3 Economics & incentives

- Use time-varying tariffs and demand response signals to align user charging behavior; include operator/aggregator revenue models for RL policies that enable V2G services. DOE VGI assessments highlight V2G value streams (ancillary services, peak shaving).

10. Novel / Unpublished contributions (you can claim)

1. **ST-G-LSTM:** A combined spatial-attention GNN + LSTM with quantile outputs for station-level probabilistic short-term forecasts (architecture and training recipe are original).
2. **C-MARL:** A CTDE MARL method with online Lagrangian multiplier updates to enforce hard distribution constraints in expectation and a fairness debt mechanism — new in the charging scheduling context.
3. **Federated TCVAE + causal explanations:** An unsupervised federated temporal CVAE with on-device causal explanation postprocessing that reduces false positives and provides human-actionable maintenance signals.
4. **Open synthetic generator:** A reproducible, seedable joint mobility-charging-grid scenario generator (configurable presets) for benchmarking forecasting, scheduling, and maintenance jointly.

(These algorithmic designs, evaluation protocols, and the synthetic generator are novel in the combination and application presented here; they form the technical core of the original contribution.)

11. Limitations and risks

- **Data access:** operator data may be restricted; results on public datasets may not fully generalize. Use of federated approaches mitigates some data sharing constraints.
- **Hardware heterogeneity:** sensor quality varies; model robustness must be validated per vendor.
- **Regulatory constraints:** V2G requires policy and tariff support in many jurisdictions.

12. Roadmap for future work

1. **Field trial:** deploy forecasting + C-MARL in a medium-sized workplace campus (50–150 chargers) and compare to rule-based baseline.
2. **V2G extension:** extend MARL to scheduled bi-directional flows and economic dispatch considering battery degradation.
3. **Edge acceleration:** quantize forecasting models for edge inferencing on local gateway devices.
4. **Adversarial robustness & cyber-security:** integrate intrusion detection for EVSE communications; explore distributed ledger for tamper-evident logs.

13. Conclusion

Integrating forecasting, constrained RL scheduling, and federated predictive maintenance can materially improve EV charger uptime, reduce grid stress, and provide fairer access to charging. The proposed ST-G-LSTM, C-MARL, and federated maintenance pipeline form a coherent, deployable suite that addresses technical, privacy, and operational challenges. We provide a reproducible experimental blueprint and recommend staged field trials to validate practical benefits.

References (selected, key sources cited in text)

1. IEA — *Global EV Outlook 2025: Electric vehicle charging*. (Summary statistics on global charging rollout). [IEA](#)
2. Caltech EV Research — *ACN-Data — A Public EV Charging Dataset*. ev.caltech.edu+1
3. Zhao, Z. et al. — *Reinforcement learning for electric vehicle charging: a review* (2024). (RL review for charging scheduling). [ScienceDirect](#)

4. U.S. Department of Energy — *Vehicle-Grid Integration Assessment Report* (2025). (VGI potentials & integration guidance). [The Department of Energy's Energy.gov](https://www.energy.gov)
5. ChargerHelp / UC Davis analysis reported in The Verge — *large study of real-world charger uptime and availability issues* (2024). (Motivates predictive maintenance). [The Verge](https://www.theverge.com)